

MongoDB Lab

Use the [NixOS Database Systems Virtual Machine](#) [../nixos-dbs-vm/] to run the code in this lab.

Demonstration Example

Run [MongoDB shell](#) [https://docs.mongodb.com/manual/mongo/] by `mongo` and submit the following MongoDB statements (except for the first code to import a dataset).

See [SQL to MongoDB Mapping Chart](#)

[https://docs.mongodb.com/manual/reference/sql-comparison/] to understand the MongoDB terminology and concepts.

Import a Dataset

```
curl -L http://media.mongodb.org/zipcodes.json | mongoimport -d test -c zipcodes
```

Query Data

[Find data in a collection](#)

[https://docs.mongodb.com/manual/reference/method/db.collection.find/] based on a query with several filters described in a type bracketing.

```
db.zipcodes.find({})
db.zipcodes.find( {
  state: "MA",
  $or: [ { pop: { $lt: 200 } }, { pop: { $gt: 10000 } } ]
} )
```

Aggregate Data

Create an aggregation pipeline

[<https://docs.mongodb.com/manual/core/aggregation-pipeline/>] to aggregate and filter data in several steps.

```
// returns all states with total population greater than 10 million
db.zipcodes.aggregate( [
  { $group: { _id: "$state", totalPop: { $sum: "$pop" } } },
  { $match: { totalPop: { $gte: 10*1000*1000 } } }
] )

// returns the average populations for cities in each state
db.zipcodes.aggregate( [
  { $group: { _id: { state: "$state", city: "$city" }, pop: {
    $sum: "$pop" } } },
  { $group: { _id: "$_id.state", avgCityPop: { $avg: "$pop" } } }
] )

// returns the smallest and largest cities by population for each
state
db.zipcodes.aggregate( [
  { $group:
    {
      _id: { state: "$state", city: "$city" },
      pop: { $sum: "$pop" }
    }
  },
  { $sort: { pop: 1 } },
  { $group:
    {
      _id : "$_id.state",
      biggestCity: { $last: "$_id.city" },
      biggestPop: { $last: "$pop" },
      smallestCity: { $first: "$_id.city" },
      smallestPop: { $first: "$pop" }
    }
  },
  // the following $project is optional, and modifies the output
  format.
  { $project:
    {
      _id: 0,
      state: "$_id",
      biggestCity: { name: "$biggestCity", pop:
        "$biggestPop" },
      smallestCity: { name: "$smallestCity", pop:
```

```
"$smallestPop" }  
    }  
  }  
] )
```

Run Map-Reduce

Create and apply map and reduce functions

[<https://docs.mongodb.com/manual/reference/method/db.collection.mapReduce/>] to process data.

```
// emit the state and the population for each city  
var mapFunction1 = function() {  
    emit(this.state, this.pop);  
};  
  
// reduce the populations array to the sum of its elements  
var reduceFunction1 = function(keyCustId, valuesPops) {  
    return Array.sum(valuesPops);  
};  
  
// calculate the average population per state  
var finalizeFunction1 = function (key, reducedVal) {  
    reducedVal.avg = reducedVal.qty/reducedVal.count;  
    return reducedVal;  
};  
  
// apply  
db.zipcodes.mapReduce(mapFunction1, reduceFunction1, {  
    out: { merge: "states_populations" },  
    query: { pop: { $gt: 10000 } },  
    finalize: finalizeFunction1  
})  
  
// show  
db.states_populations.find({})
```